

CSCE 435 Group project

0. Group number: 26

1. Group members:

1. Caroline Jia

- Username: cjia2003
- Algorithm: Radix Sort

2. Griffin Beaudreau

- Username: CyberGriffin
- Algorithm: Column Sort

3. Kaitlyn Griffin

- Username: kaitlyngrif
- Algorithm: Sample Sort

4. Samuel Bush

- Username: SamShrubo
- Algorithm: Merge Sort

5. Zhongyou Wu

- Username: ZhongyouWuTAMU
- Algorithm: Bitonic Sort

1a. Team Communication:

We will be using Discord for our team communications.

2. Project topic: Parallel Sorting Algorithms

2a. Brief project description (what algorithms will you be comparing and on what architectures)

- Bitonic Sort:
- Sample Sort: Will be implemented using MPI on the Grace cluster. The initial large problem array will split into multiple sub-arrays to be distributed across Grace's nodes and processors.
- Merge Sort: Implement using MPI on the Grace cluster, split the initial array into multiple sub-arrays to distribute across the network of nodes and processors
- Radix Sort: Implemented using MPI on Grace cluster. The initial array will be split into multiple smaller arrays across the nodes and processors, will be using least significant digit version
- Column Sort: A parallel sorting algorithm that is well suited for sorting data arranged in a 2D grid. The matrix is sorted column-wise, transposed, and sorted again row-wise. This process is repeated until the matrix is sorted. This algorithm will be implemented using MPI on the Grace cluster.

2b. Pseudocode for each parallel algorithm

- For MPI programs, include MPI calls you will use to coordinate between processes

****---Merge Sort Pseudocode---****

```c++

# Full sorting algorithm

def MergeSort(Array, arraySize) {

    MPI\_Init(arguments to set up mpi)

    MPI\_Comm\_rank(MPI\_COMM\_WORLD, rank)    # set rank to this process rank

    MPI\_Comm\_size(MPI\_COMM\_WORLD, numProcs) # get number of processes

    # divide array into local sub-arrays for each process to sort individually

    if arraySize % numProcs != 0:

        # round up for array size and any unfilled space in the sub-array can be accounted for as null

        localArraySize = ceiling(arraySize / numProcs)

    else:

        localArraySize = arraySize / numProcs

```

make local array

arrayOffset = rank * localArraySize

localArray = array[localArraySize] # create buffer for receiving scattered data

scatter to all processes

MPI_Scatter(Array, localArraySize, MPI_INT, localArray, localArraySize, MPI_INT, root=0,
MPI_COMM_WORLD)

each process sorts locally using quicksort

localQuickSort(localArray, localArraySize)

begin merging with neighbor processes

step = 1

while step < num_procs:

 # combine even and odd processes in the even process

 if (rank % (2 * step) == 0):

 if (rank + step < num_procs):

 # get sorted array from neighbor process

 receivedSize = localArraySize * step

 receivedArray = new array[receivedSize]

 MPI_Recv(receivedArray, receivedSize, MPI_INT, rank + step, 0, MPI_COMM_WORLD,
MPI_STATUS_IGNORE)

 # Merge local array and received array using helper function defined below

 localArray = Merge(localArray, localArraySize, receivedArray, receivedSize)

 # Update local size after merge

 localArraySize = localArraySize + receivedSize

```

```

elif (rank % step == 0):
 # Send local array to neighboring process
 MPI_Send(localArray, localArraySize, MPI_INT, rank - step, 0, MPI_COMM_WORLD)

 # double step size each iteration
 step = step * 2

only main process gets final sorted array
if rank == 0
 sortedArray = array[arraySize]
else
 sortedArray = null

from all scattered processes gather into final sorted array
MPI_Gather(localArray, localArraySize, MPI_INT, sortedArray, localArraySize, MPI_INT, root=0,
MPI_COMM_WORLD)

only main can display / output the sorted array
if rank == 0
 Display(sortedArray)

call in every process
MPI_Finalize()
}

helper function merges 2 already sorted arrays into 1
def Merge(Array1, array1Size, Array2, array2Size) {
 mergedArray = array[array1Size + array2Size]
 i = 0

```

```
j = 0
```

```
k = 0
```

```
compare each element until completed 1 array
```

```
while (i < array1Size) and (j < array2Size)
```

```
 if Array1[i] < Array2[j]
```

```
 mergedArray[k] = Array1[i]
```

```
 i++
```

```
 else
```

```
 mergedArray[k] = Array2[j]
```

```
 j++
```

```
 k++
```

```
if unread elements in Array1 copy them to the end of mergedArray
```

```
while i < array1Size
```

```
 mergedArray[k] = Array1[i]
```

```
 i++
```

```
 k++
```

```
if unread elements in Array2 copy them to the end of mergedArray
```

```
while j < array2Size
```

```
 mergedArray[k] = Array2[j]
```

```
 j++
```

```
 k++
```

```
return mergedArray
```

```
}
```

```
...
```

**\*\*---Sample Sort Pseudocode---\*\***

```\n

sample sort pseudocode here

```\n

**\*\*---Radix Sort Pseudocode---\*\***

```\n

MPI_Init()

MPI_Comm_rank(comm, rank);

MPI_Comm_size(comm, size);

// Scatter the array across processes

if (rank == 0) {

int *arr = generate_input_array(N);

// Scatter the array to all processes

MPI_Scatter(variables);

} else {

// Other processes

MPI_Scatter(variables);

}

// Perform radix sort on the local portion of the array

while i < max_digits {

// Perform counting sort at current digit

int local_count = counting_sort_by_digit(local_arr, local_size, digit_pos, base);

// Gather global counts from other processes

```

MPI_Allgather(local_count, base, MPI_INT, global_count, base, MPI_INT, comm);

// Compute prefix sums on global counts to determine offsets
int prefix_sum = compute_prefix_sums(global_count, base);

// Redistribute elements based on the computed prefix sums
int sorted_local_arr = redistribute_elements(local_arr, local_size, digit_pos, prefix_sum, base);

// Replace local array with the newly sorted portion
local_arr = sorted_local_arr;
}

// Gather the locally sorted arrays back into the root process
if (rank == 0) {
    MPI_Gather(local_arr, local_size, MPI_INT, sorted_arr, local_size, MPI_INT, 0,
MPI_COMM_WORLD);
} else {
    MPI_Gather(local_arr, local_size, MPI_INT, NULL, local_size, MPI_INT, 0, MPI_COMM_WORLD);
}

MPI_Finalize();

...

```

****---Column Sort Pseudocode---****

...

Steps:

1: Arrange data in a matrix with r rows and c columns, where r is the number of processors and c is the number of items per process.

2: Sort each column using a sequential sorting algorithm (may change algorithm depending on input size).

3: Transpose the matrix.

4: Sort each row independently

5: Sort each column again.

6: Tranpose the matrix.

7: Final column sort.

```

```c++

/*

* Include MPI header

* Include Caliper header

* Include any additional headers

* Define Constants (MASTER)

*/

int main(int argc, char *argv[]) {

CALI_CXX_MARK_FUNCTION;

if args invalid return 0;

// Set up MPI environment (needs to include arguments)

MPI_INIT();

MPI_Comm_rank()

MPI_Comm_size()

// Initialize variables

int N = atoi(argv[1])

int P = num processors;

int rows, N_padded, padding_size;


```
// To handle cases where the total number of data elements isn't divisible by the number of processors
```

```
// Example: N = 10 and P = 4:
```

```
// rows = (10 + 4 - 1) / 4 = 3
```

```
// N_padded = 3 * 4 = 12 (what it will be when padded)
```

```
// padding_size = 12 - 10 = 2 (num of padding elements)
```

```
rows = (N + P - 1) / P;
```

```
N_padded = rows * P;
```

```
padding_size = N_padded - N;
```

```
// Local data for each processor
```

```
int *local_array = (int*)malloc(rows * sizeof(int));
```

```
if (taskid == MASTER) {
```

```
    int *global_data = (int*)malloc(N_padded * sizeof(int));
```

```
    /* Add actual data to global_data */
```

```
    /* Add padded data */
```

```
    for (int i = N; i < N_padded; ++i) {
```

```
        global_data[i] = INT_MAX;
```

```
    }
```

```
    for (int p = 0; p < P; ++p) {
```

```
        if (p == MASTER) {
```

```
            // copy local data to global data
```

```
        } else {
```

```
            // MPI_Send
```

```
        }
```

```

    }
    free(global_data)
} else {
    // MPI_Recv
}

/* local column sort, marked with caliper */

/* transpose, marked with caliper */

/* lcaol row sort, marked with caliper */

/* transpose, marked with caliper */

/* local column sort, marked with caliper */

if (taskid == MASTER) {
    int *sorted_data = (int*)malloc(N * sizeof(int));
    // copy local_data into sorted_data

    foreach p = 1; p < column,
MPI_Recv(sorted_data[p*rows],rows,MPI_INIT,p,0,MPI_COMM_WORLD,MPI_STATUS_IGNORE);

    // Verify that array is sorted

    free(sorted_data);
} else {
    MPI_SNED(local_data, rows, MPI_INT, MASTER, 0, MPI_COMM_WORLD);
}

```

```

free(local_data);

mgr.flush();

MPI_Finalize();

return 0;
}
...

```

2c. Evaluation plan - what and how will you measure and compare

- Evaluating with multiple process counts, the total process count should always be a power of 2 (2^n processors):
 - Processor count: 2, 4, 8, 16, 32, 64, 128, 256, 512, 1024
- Using Caliper + thicket to calculate the total time taken, time per process, etc, for each algorithm with the same inputs and processor count across each
 - This method can allow us to determine which algorithms are fastest in what input context
- Adjust the following in each evaluation case to test each algorithm:
 - Input sizes, Input types:
 - Input sizes: 2^{16} , 2^{18} , 2^{20} , 2^{22} , 2^{24} , 2^{26} , 2^{28}
 - Input types: Sorted, Random, Reverse sorted, 1% perturbed
 - Strong scaling (same problem size, increase the number of processors/nodes)
 - Weak scaling (increase problem size, increase the number of processors)